

This is a pdf version of <https://git.tu-berlin.de/lis-public/ai-student-workspace/-/blob/main/01/README.md>

To obtain the code for this assignment, you will need to fetch and pull new commits from `git@git.tu-berlin.de:lis-public/ai-student-workspace.git`

As always, only modify the file `solution_??.py`. And even in `solution_??.py`, only modify what the functions do – don't change the function's names. Of course, you can import additional libraries, and add files with own code to import.

You can run tests by changing directory to the task folder `??`, and then simply typing `python3 -m pytest`. If you haven't yet, you will need to install pytest first: `sudo apt install python3-pytest`.

# Assignment 1: Single-Step and Discrete Decision Making

---

## 1.1: Single Step, Discrete, Full knowledge

The following information is available:

- There are  $A$  possible choices.
- Depending on your choice  $a \in 0, 1, \dots, A - 1$ , the outcome  $x \in \mathbb{R}$  will be Gaussian distributed with mean  $\mu_a$  and standard deviation  $\sigma_a$ , in other words:  $P(x|a) = N(x; \mu_a, \sigma_a)$ . We have full knowledge of the outcome of our actions, meaning that all  $\mu_a$  and  $\sigma_a$  are known.

For this part of the exercise, you will need to modify the following three functions in `solution_01.py`:

1. `maximize_mean(mu_vector, sigma_vector)`
2. `maximize_ucb(mu_vector, sigma_vector)`
3. `maximize_utility(mu_vector, sigma_vector, utility_function)`

The inputs are as follows:

- `mu_vector` is a `numpy.array` vector of length  $A$ . It contains the  $\mu_a$  for each of the  $A$  possible choices.
- `sigma_vector` is a `numpy.array` vector of length  $A$ . It contains the  $\sigma_a$  for each of the  $A$  possible choices.
- `utility_function` is only input to the last function, not to the others. `utility_function` is a function  $R : \mathbb{R} \rightarrow \mathbb{R}$ . The function object `utility_function` can be called using any real number  $x \in \mathbb{R}$ , and will return another real number  $y$ .

Each of these functions outputs a decision  $a \in 1, 2, \dots, A$ . Please modify these functions as follows:

### 1.1.1: Maximum Mean (1 Point)

Modify `maximize_mean(mu_vector, sigma_vector)` so that it outputs the decision  $a$  that maximizes the mean outcome  $\int P(x|a) x \, dx = \mu_a$ . (Yes, this is trivial to implement.)

### 1.1.2: Maximum UCB (1 Point)

Modify `maximize_ucb(mu_vector, sigma_vector)` so that it outputs the decision  $a$  that maximizes the maximum upper confidence bound  $\mu_a + \beta\sigma_a$  for  $\beta = 2$ .

### 1.1.3: Maximum Utility (1 Point)

Modify `maximize_utility(mu_vector, sigma_vector, utility_function)` so that it outputs the decision  $a$  that maximizes the expected utility with respect to an arbitrary utility function  $R : \mathbb{R} \rightarrow \mathbb{R}$ . In other words, it must output the decision  $a$  that maximizes

$$\int P(x|a) R(x) \, dx \quad . \quad (1)$$

The utility function  $R$  is given to `maximize_utility` as the function object `utility_function`, as described above.

We assume no knowledge on  $R$ , we can only query it for different values of  $x$ . That means that in order to evaluate the integral in eq. (1), we need to use numerical methods like Monte-Carlo sampling. We assume a limited budget, i.e. the function  $R$  can only be sampled  $n = 1000$  times.

## 1.2: Single Step, Discrete, Simulation-based (aka Bandits) (2 Points)

The following information is available:

- There are  $A$  possible choices.
- Depending on your choice  $a \in 0, 1, \dots, A - 1$ , we know that the outcome  $x \in [0, 1]$  will be strictly in the interval  $[0, 1]$ . However, in contrast to 1.1, the distribution  $P(x|a)$  is not known (obviously, it is not Gaussian!).
- We have access to a stochastic simulation that we can access  $n$  times. Each time, the agent queries an action  $a_i$ , and the simulation returns an outcome  $x_i$ , where, again,  $x_i \in [0, 1]$ . Put differently, the agent makes  $n$  data-collecting decisions to collect data  $D = (a_i, x_i)_{i=1}^n$ .

For this part of the exercise, you will need to modify the following function in `solution_01.py`:

```
maximize_mean_using_simulator(simulator, A, n)
```

The inputs are as follows:

- `simulator` is a function that simulates the outcome of the actions. It returns the outcome  $x = \text{simulator}(a)$  depending on the action  $a \in 0, 1, \dots, A - 1$ .
- `A` is the number of actions that are possible. Your method must output an integer  $a \in 0, 1, \dots, A - 1$ .
- `n` is the budget for querying `simulator`. If you query `simulator` more often, it will raise an exception (error).

Modify `maximize_mean_using_simulator` so that it returns the action  $a \in 0, 1, \dots, A - 1$  that maximizes the expected mean outcome  $\int P(x|a) x \, dx$ . This is analogous to 1.1.1, but this time you don't know  $P$ .